

LinearRegression

In [1]:

```
import numpy as np
import pandas as pd
```

data collection

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as pp
import seaborn as sb
```

In [3]:

```
df = pd.read_csv(r"C:\Users\user\Desktop\5_Instagram data.csv")  
df
```

Out[3]:

	Impressions	From Home	From Hashtags	From Explore	From Other	Saves	Comments	Shares	Likes	Profile Visits
0	3920	2586	1028	619	56	98	9	5	162	36
1	5394	2727	1838	1174	78	194	7	14	224	48
2	4021	2085	1188	0	533	41	11	1	131	62
3	4528	2700	621	932	73	172	10	7	213	29
4	2518	1704	255	279	37	96	5	4	123	8
...
114	13700	5185	3041	5352	77	573	2	38	373	73
115	5731	1923	1368	2266	65	135	4	1	148	20
116	4139	1133	1538	1367	33	36	0	1	92	34
117	32695	11815	3147	17414	170	1095	2	75	549	148
118	36919	13473	4176	16444	2547	653	5	26	443	61

119 rows × 13 columns


first 10 rows

In [4]:

```
df.head(10)
```

Out[4]:

	Impressions	From Home	From Hashtags	From Explore	From Other	Saves	Comments	Shares	Likes	Profile Visits
0	3920	2586	1028	619	56	98	9	5	162	35
1	5394	2727	1838	1174	78	194	7	14	224	48
2	4021	2085	1188	0	533	41	11	1	131	62
3	4528	2700	621	932	73	172	10	7	213	23
4	2518	1704	255	279	37	96	5	4	123	8
5	3884	2046	1214	329	43	74	7	10	144	9
6	2621	1543	599	333	25	22	5	1	76	26
7	3541	2071	628	500	60	135	4	9	124	12
8	3749	2384	857	248	49	155	6	8	159	36
9	4115	2609	1104	178	46	122	6	3	191	31



data cleaning

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 119 entries, 0 to 118
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Impressions            119 non-null    int64
1   From Home              119 non-null    int64
2   From Hashtags          119 non-null    int64
3   From Explore           119 non-null    int64
4   From Other             119 non-null    int64
5   Saves                  119 non-null    int64
6   Comments               119 non-null    int64
7   Shares                 119 non-null    int64
8   Likes                  119 non-null    int64
9   Profile Visits         119 non-null    int64
10  Follows                119 non-null    int64
11  Caption                119 non-null    object
12  Hashtags               119 non-null    object
dtypes: int64(11), object(2)
memory usage: 12.2+ KB
```

In [6]:

```
df.describe()
```

Out[6]:

	Impressions	From Home	From Hashtags	From Explore	From Other	Saves	Comments
count	119.000000	119.000000	119.000000	119.000000	119.000000	119.000000	119.000000
mean	5703.991597	2475.789916	1887.512605	1078.100840	171.092437	153.310924	107.000000
std	4843.780105	1489.386348	1884.361443	2613.026132	289.431031	156.317731	107.000000
min	1941.000000	1133.000000	116.000000	0.000000	9.000000	22.000000	0.000000
25%	3467.000000	1945.000000	726.000000	157.500000	38.000000	65.000000	0.000000
50%	4289.000000	2207.000000	1278.000000	326.000000	74.000000	109.000000	0.000000
75%	6138.000000	2602.500000	2363.500000	689.500000	196.000000	169.000000	0.000000
max	36919.000000	13473.000000	11817.000000	17414.000000	2547.000000	1095.000000	119.000000

In [7]:

```
df.columns
```

Out[7]:

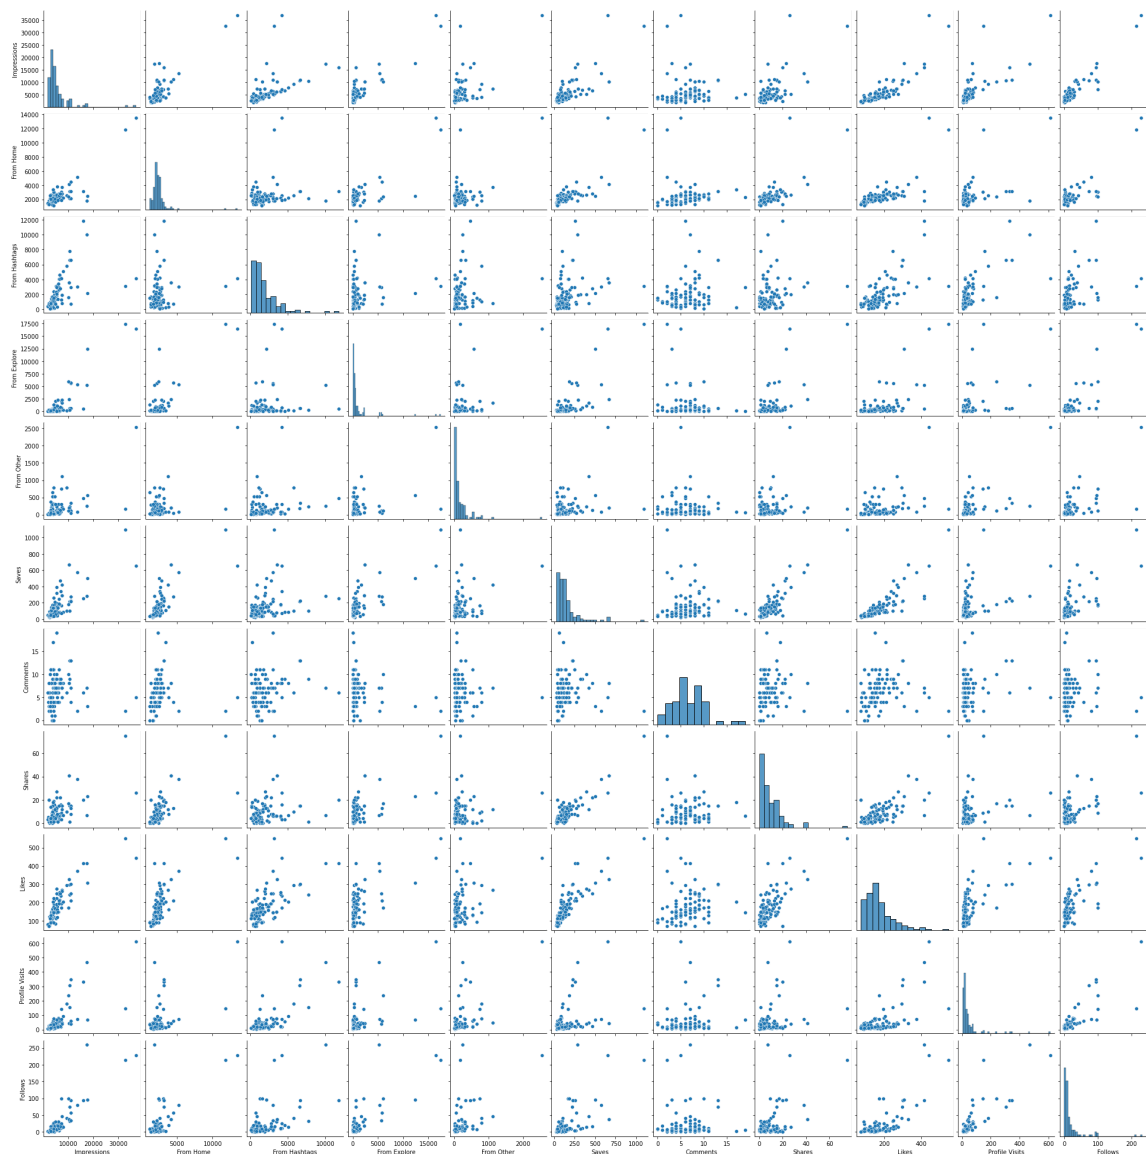
```
Index(['Impressions', 'From Home', 'From Hashtags', 'From Explore',  
      'From Other', 'Saves', 'Comments', 'Shares', 'Likes', 'Profile Visi  
ts',  
      'Follows', 'Caption', 'Hashtags'],  
      dtype='object')
```

In [8]:

```
sb.pairplot(df)
```

Out[8]:

```
<seaborn.axisgrid.PairGrid at 0x1d255e374c0>
```



In [9]:


```
sb.distplot(df["Shares"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

<AxesSubplot:xlabel='Shares', ylabel='Density'>



In [10]:


```
sb.distplot(df["Shares"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[10]:

<AxesSubplot:xlabel='Shares', ylabel='Density'>



In [11]:

```
df1=df[['Impressions', 'From Home', 'From Hashtags', 'From Explore',
        'From Other', 'Saves', 'Comments', 'Shares', 'Likes', 'Profile Visits',
        'Follows']]
df1
```

Out[11]:


	Impressions	From Home	From Hashtags	From Explore	From Other	Saves	Comments	Shares	Likes	Profile Visits	Follows
0	3920	2586	1028	619	56	98	9	5	162	35	2
1	5394	2727	1838	1174	78	194	7	14	224	48	10
2	4021	2085	1188	0	533	41	11	1	131	62	12
3	4528	2700	621	932	73	172	10	7	213	23	8
4	2518	1704	255	279	37	96	5	4	123	8	0
...
114	13700	5185	3041	5352	77	573	2	38	373	73	80
115	5731	1923	1368	2266	65	135	4	1	148	20	18
116	4139	1133	1538	1367	33	36	0	1	92	34	10
117	32695	11815	3147	17414	170	1095	2	75	549	148	214

In [12]:

```
sb.heatmap(df1.corr())
```

Out[12]:

<AxesSubplot:>



model building

In [13]:

```
x = df1[['Impressions', 'From Home', 'From Hashtags', 'From Explore',
        'From Other', 'Saves', 'Comments', 'Shares', 'Likes', 'Profile Visits',
        'Follows']]
y = df1['Shares']
```

In [14]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear regression

In [15]:

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(x_train,y_train)
```

Out[15]:

LinearRegression()

In [16]:

```
print(lr.intercept_)
```

-9.592326932761353e-14

In [17]:

```
coef = pd.DataFrame(lr.coef_,x.columns,columns=['Co_efficient'])
coef
```

Out[17]:

	Co_efficient
Impressions	6.557359e-17
From Home	3.659749e-18
From Hashtags	-1.044350e-16
From Explore	-6.584377e-17
From Other	2.000772e-16
Saves	-8.700698e-17
Comments	2.761500e-16
Shares	1.000000e+00
Likes	-5.278061e-17
Profile Visits	-4.977442e-16
Follows	1.607783e-16

In [18]:

```
print(lr.score(x_test,y_test))
```


1.0

In [19]:

```
prediction = lr.predict(x_test)
pp.scatter(y_test,prediction)
```

Out[19]:

<matplotlib.collections.PathCollection at 0x1d2603a4cd0>



lasso and ridge regression

In [20]:

```
lr.score(x_test,y_test)
```

Out[20]:

1.0

In [21]:

```
lr.score(x_train,y_train)
```

Out[21]:

1.0

In [22]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [23]:

```
r = Ridge(alpha=10)
r.fit(x_train,y_train)
r.score(x_test,y_test)
r.score(x_train,y_train)
```

Out[23]:

0.9999946903983258

In [24]:

```
l = Lasso(alpha=10)
l.fit(x_train,y_train)
l.score(x_test,y_test)
l.score(x_train,y_train)
```

Out[24]:

0.9674494232346373

elasticnet

In [25]:

```
from sklearn.linear_model import ElasticNet
e = ElasticNet()
e.fit(x_train,y_train)
```

Out[25]:

ElasticNet()

In [26]:

```
print(e.coef_)
```

```
[ 4.65781471e-06  1.24855628e-04  0.00000000e+00 -2.70638912e-05
 -2.12048393e-04  1.90826793e-03  0.00000000e+00  9.57608056e-01
 -0.00000000e+00 -3.02663893e-05 -0.00000000e+00]
```

In [27]:

```
print(e.intercept_)
```

-0.14962262072909205

In [28]:

```
predictions = e.predict(x_test)
predictions
```

Out[28]:

```
array([ 6.22575062,  3.0321759 ,  3.29083321,  6.00451146,  0.09114064,
        3.28541562, 10.83768675,  2.29562028,  8.25006063,  2.29514008,
       15.5747305 ,  1.25667766, 16.6262233 ,  3.27029655, 13.82789928,
        5.26003122,  6.00451146,  3.89309877,  9.82292084, 22.76903267,
        8.34076266, 12.90713554,  2.06863276,  8.02274936, 21.84623472,
        1.04271171, 13.6948035 , 26.84480889, 10.82069932, 14.80039615,
        1.11225466,  0.12878671,  1.03972728, 15.03628293,  1.96900648,
        4.17065627])
```

In [29]:

```
print(e.score(x_test,y_test))
```

0.9987041517146851

In [30]:

```
from sklearn import metrics
```

mean absolute error

In [31]:

```
print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,predictions))
```

Mean Absolute Error: 0.19387974272182104

mean squared error

In [32]:

```
print("Mean Squared Error:", metrics.mean_squared_error(y_test,predictions))
```

Mean Squared Error: 0.062196718163370884

root mean squared error

In [33]:

```
print("Root Mean Squared Error",np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

Root Mean Squared Error 0.2493926986969965

In []: