

LinearRegression

In [1]:

```
import numpy as np
import pandas as pd
```

data collection

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as pp
import seaborn as sb
```

In [3]:

```
df = pd.read_csv(r"C:\Users\user\Desktop\11_winequality-red.csv")
df
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	

1599 rows × 12 columns

first 10 rows

In [4]:

```
df.head(10)
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alco
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	1
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	1

data cleaning

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [6]:

```
df.describe()
```

Out[6]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.406864
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.847046
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

In [7]:

```
df.columns
```

Out[7]:

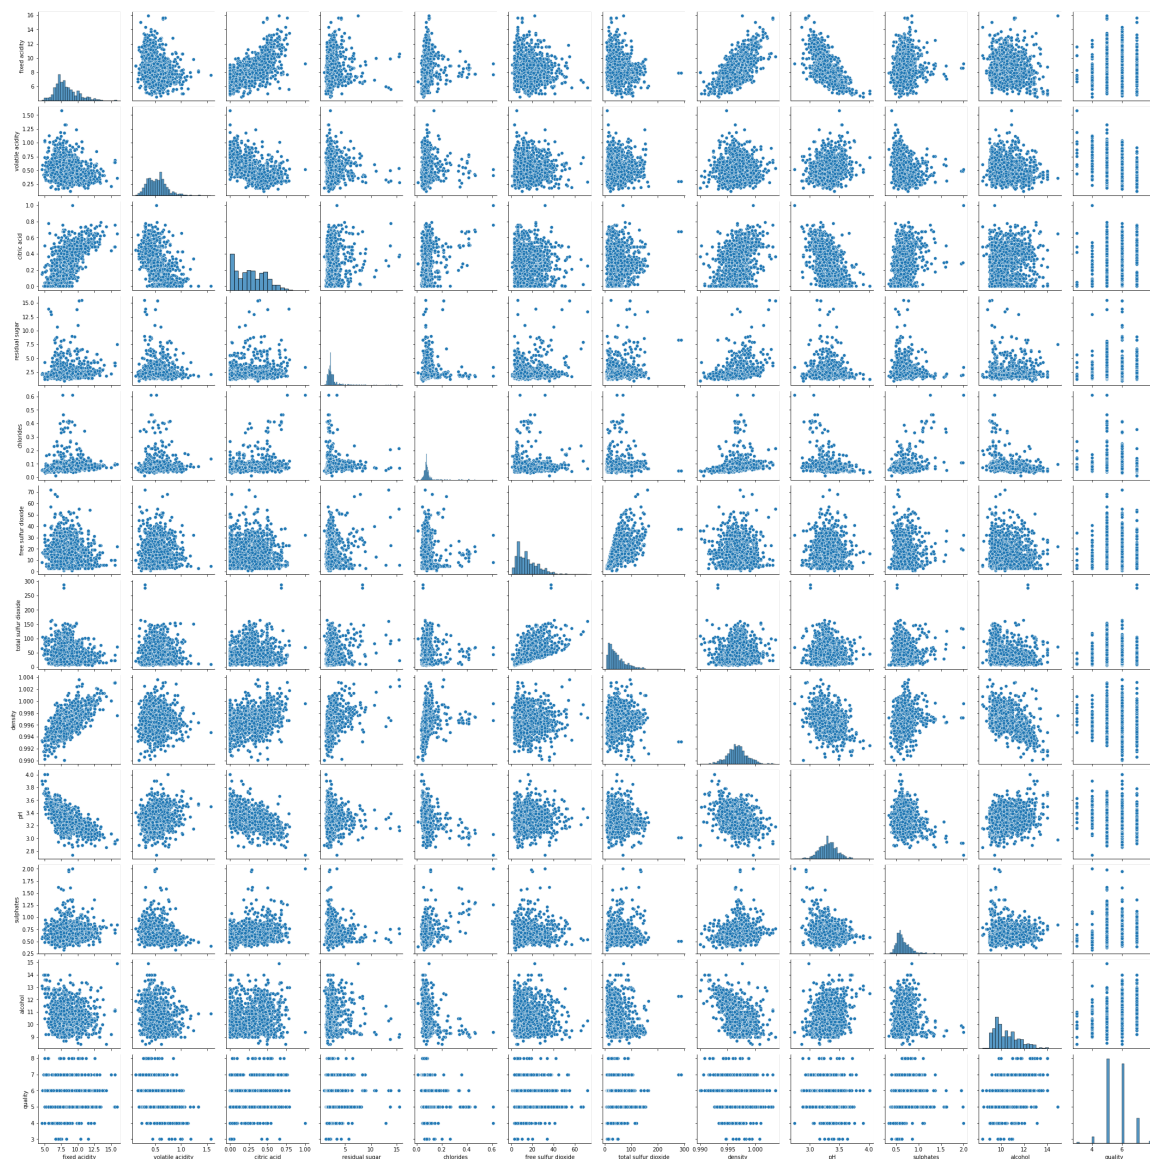
```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
      'pH', 'sulphates', 'alcohol', 'quality'],  
      dtype='object')
```

In [8]:

```
sb.pairplot(df)
```

Out[8]:

<seaborn.axisgrid.PairGrid at 0x209a025df70>



In [9]:


```
sb.distplot(df["density"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

<AxesSubplot:xlabel='density', ylabel='Density'>



In [10]:


```
sb.distplot(df["density"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[10]:

<AxesSubplot:xlabel='density', ylabel='Density'>



In [11]:

```
df1=df[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
        'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
        'pH', 'sulphates', 'alcohol', 'quality']]
df1
```

Out[11]:


	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.0	0.540	0.10	2.0	0.070	30.0	40.0	0.99574	3.40	0.75	11.0	

In [12]:

```
sb.heatmap(df1.corr())
```

Out[12]:

<AxesSubplot:>



model building

In [15]:

```
x = df1[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
        'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
        'pH', 'sulphates', 'alcohol', 'quality']]  
y = df1['density']
```

In [16]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear regression

In [17]:

```
from sklearn.linear_model import LinearRegression  
  
lr = LinearRegression()  
lr.fit(x_train,y_train)
```

Out[17]:

LinearRegression()

In [18]:

```
print(lr.intercept_)
```

4.440892098500626e-15

In [19]:

```
coef = pd.DataFrame(lr.coef_,x.columns,columns=['Co_efficient'])  
coef
```

Out[19]:

	Co_efficient
fixed acidity	4.638736e-17
volatile acidity	-3.741660e-16
citric acid	-4.503987e-16
residual sugar	-6.525389e-18
chlorides	5.598087e-15
free sulfur dioxide	4.208415e-18
total sulfur dioxide	-1.407067e-19
density	1.000000e+00
pH	3.803632e-16
sulphates	-5.457805e-16
alcohol	3.796431e-17
quality	3.061217e-17

In [20]:

```
print(lr.score(x_test,y_test))
```


1.0

In [21]:

```
prediction = lr.predict(x_test)  
pp.scatter(y_test,prediction)
```

Out[21]:

<matplotlib.collections.PathCollection at 0x209a9d36790>



lasso and ridge regression

In [22]:

```
lr.score(x_test,y_test)
```

Out[22]:

1.0

In [23]:

```
lr.score(x_train,y_train)
```

Out[23]:

1.0

In [24]:

```
from sklearn.linear_model import Ridge,Lasso
```


In [25]:

```
r = Ridge(alpha=10)
r.fit(x_train,y_train)
r.score(x_test,y_test)
r.score(x_train,y_train)
```

Out[25]:

0.8213534011567556

In [26]:

```
l = Lasso(alpha=10)
l.fit(x_train,y_train)
l.score(x_test,y_test)
l.score(x_train,y_train)
```

Out[26]:

0.0

elasticnet

In [27]:

```
from sklearn.linear_model import ElasticNet
e = ElasticNet()
e.fit(x_train,y_train)
```

Out[27]:

ElasticNet()

In [28]:

```
print(e.coef_)
```

[0. 0. 0. 0. 0. -0. 0. 0. -0. 0. -0. -0.]

In [29]:

```
print(e.intercept_)
```

0.9967834405719392

In [30]:

```
predictions = e.predict(x_test)
predictions
```

Out[30]:

In [31]:

-0.004107190678244670799678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344.

In [32]:

```
from sklearn import
```

mean absolute error 0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344

In [33]:

```
print("Mean Absolute Error:", metrics.mean_absolute_error(y_test, pr
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
```

Mean Ab₉₀ 29678344 or 0.096783445205

mean squared error

```
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
In [34]: 0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
print("Mean Squared Error:", metrics.mean_squared_error(y_test, predictions))
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
Mean Squared Error: 3.86523188888592e-06
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
root mean squared error
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
In [35]: 0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
print("Root Mean Squared Error", np.sqrt(metrics.mean_squared_error(y_test, predictions)))
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
Root Mean Squared Error: 0.0019657840350799678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
In [ ]: 0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
0.99678344, 0.99678344, 0.99678344, 0.99678344, 0.99678344,
```