

LinearRegression

In [1]:

```
import numpy as np
import pandas as pd
```

data collection

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as pp
import seaborn as sb
```

In [3]:

```
df = pd.read_csv(r"C:\Users\user\Desktop\14_Iris.csv")
df
```

Out[3]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

first 10 rows

In [4]:

```
df.head(10)
```

Out[4]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
5	6	5.4	3.9	1.7	0.4	Iris-setosa
6	7	4.6	3.4	1.4	0.3	Iris-setosa
7	8	5.0	3.4	1.5	0.2	Iris-setosa
8	9	4.4	2.9	1.4	0.2	Iris-setosa
9	10	4.9	3.1	1.5	0.1	Iris-setosa

data cleaning

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id               150 non-null   int64
1   SepalLengthCm   150 non-null   float64
2   SepalWidthCm    150 non-null   float64
3   PetalLengthCm   150 non-null   float64
4   PetalWidthCm    150 non-null   float64
5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

In [6]:

```
df.describe()
```

Out[6]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

In [7]:

```
df.columns
```

Out[7]:

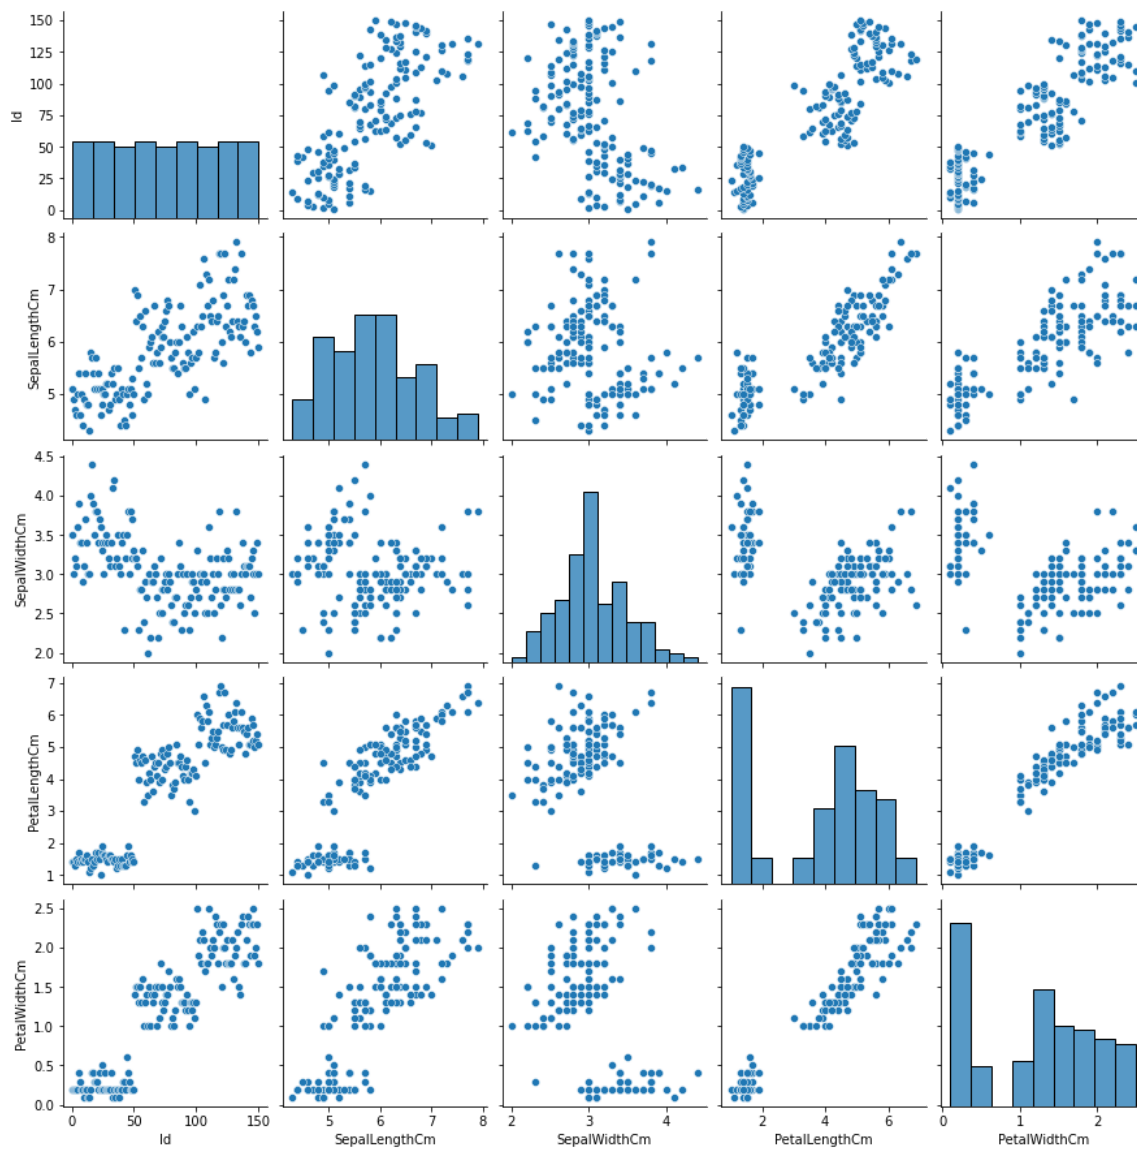
```
Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
      'Species'],  
      dtype='object')
```

In [8]:

```
sb.pairplot(df)
```

Out[8]:

<seaborn.axisgrid.PairGrid at 0x2e168aeb730>



In [9]:


```
sb.distplot(df["SepalLengthCm"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

<AxesSubplot:xlabel='SepalLengthCm', ylabel='Density'>



In [10]:


```
sb.distplot(df["SepalLengthCm"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[10]:

<AxesSubplot:xlabel='SepalLengthCm', ylabel='Density'>



In [11]:

```
df1=df[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
df1
```

Out[11]:


	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	1	5.1	3.5	1.4	0.2
1	2	4.9	3.0	1.4	0.2
2	3	4.7	3.2	1.3	0.2
3	4	4.6	3.1	1.5	0.2
4	5	5.0	3.6	1.4	0.2
...
145	146	6.7	3.0	5.2	2.3
146	147	6.3	2.5	5.0	1.9
147	148	6.5	3.0	5.2	2.0
148	149	6.2	3.4	5.4	2.3

In [12]:

```
sb.heatmap(df1.corr())
```

Out[12]:

<AxesSubplot:>



model building

In [13]:

```
x = df1[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
y = df1['SepalLengthCm']
```

In [14]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear regression

In [15]:

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(x_train,y_train)
```

Out[15]:

LinearRegression()

In [16]:

```
print(lr.intercept_)
```

-5.3290705182007514e-14

In [17]:

```
coef = pd.DataFrame(lr.coef_,x.columns,columns=['Co_efficient'])
coef
```

Out[17]:

	Co_efficient
Id	6.673969e-16
SepalLengthCm	1.000000e+00
SepalWidthCm	1.364839e-16
PetalLengthCm	1.520633e-16
PetalWidthCm	-1.359544e-16

In [18]:

```
print(lr.score(x_test,y_test))
```


1.0

In [19]:

```
prediction = lr.predict(x_test)
pp.scatter(y_test, prediction)
```

Out[19]:

<matplotlib.collections.PathCollection at 0x2e16bbc5eb0>



lasso and ridge regression

In [20]:

```
lr.score(x_test, y_test)
```

Out[20]:

1.0

In [21]:

```
lr.score(x_train, y_train)
```

Out[21]:

1.0

In [22]:

```
from sklearn.linear_model import Ridge, Lasso
```

In [23]:

```
r = Ridge(alpha=10)
r.fit(x_train, y_train)
r.score(x_test, y_test)
r.score(x_train, y_train)
```

Out[23]:

0.972119849970007

In [24]:

```
l = Lasso(alpha=10)
l.fit(x_train,y_train)
l.score(x_test,y_test)
l.score(x_train,y_train)
```

Out[24]:

0.46363770678036564

elasticnet

In [25]:

```
from sklearn.linear_model import ElasticNet
e = ElasticNet()
e.fit(x_train,y_train)
```

Out[25]:

ElasticNet()

In [26]:

```
print(e.coef_)
```

[0.01305457 0. 0. 0. 0.]

In [27]:

```
print(e.intercept_)
```

4.855534100983689

In [28]:

```
predictions = e.predict(x_test)
predictions
```

Out[28]:

```
array([5.23411673, 5.76935424, 5.87379083, 5.73019052, 5.79546339,
        6.60484694, 5.66491765, 5.20800758, 5.78240881, 5.32549874,
        6.7223381 , 6.34375547, 5.07746185, 5.45604448, 6.01739114,
        5.2993896 , 6.46124663, 5.71713595, 5.96517284, 5.14273472,
        5.0513527 , 5.91295455, 5.42993533, 6.3307009 , 5.6910268 ,
        4.86858867, 6.68317438, 5.11662557, 5.83462711, 5.53437192,
        5.41688076, 6.39597376, 5.74324509, 6.12182772, 6.23931888,
        6.35681004, 6.40902834, 5.48215363, 5.01218898, 6.44819206,
        5.40382618, 6.53957407, 4.90775239, 5.93906369, 5.61269936])
```

In [29]:

```
print(e.score(x_test,y_test))
```

0.4518273528097345

In [30]:

```
from sklearn import metrics
```

mean absolute error

In [31]:

```
print("Mean Absolute Error:", metrics.mean_absolute_error(y_test, predictions))
```

Mean Absolute Error: 0.5148158350621056

mean squared error

In [32]:

```
print("Mean Squared Error:", metrics.mean_squared_error(y_test, predictions))
```

Mean Squared Error: 0.4041588949407736

root mean squared error

In [33]:

```
print("Root Mean Squared Error", np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

Root Mean Squared Error 0.6357349250597875

In []: