LinearRegression

In [1]:

```
import numpy as np
import pandas as pd
```

data collection

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as pp
import seaborn as sb
```

In [3]:

```
df = pd.read_csv(r"C:\Users\user\Desktop\8_BreastCancerPrediction.csv")
df
```

Out[3]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothne
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	
...	...	...	...	...	...	...	
564	926424	M	21.56	22.39	142.00	1479.0	
565	926682	M	20.13	28.25	131.20	1261.0	
566	926954	M	16.60	28.08	108.30	858.1	
567	927241	M	20.60	29.33	140.10	1265.0	
568	92751	B	7.76	24.54	47.92	181.0	

569 rows × 33 columns

first 10 rows

In [4]:

```
df.head(10)
```

Out[4]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	
5	843786	M	12.45	15.70	82.57	477.1	
6	844359	M	18.25	19.98	119.60	1040.0	
7	84458202	M	13.71	20.83	90.20	577.9	
8	844981	M	13.00	21.82	87.50	519.8	
9	84501001	M	12.46	24.04	83.97	475.9	

10 rows × 33 columns

data cleaning

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
 #   Column                   Non-Null Count  Dtype  
---  ------                   --------------  -----  
 0   id                       569 non-null    int64  
 1   diagnosis                569 non-null    object 
 2   radius_mean              569 non-null    float64
 3   texture_mean             569 non-null    float64
 4   perimeter_mean           569 non-null    float64
 5   area_mean                569 non-null    float64
 6   smoothness_mean          569 non-null    float64
 7   compactness_mean         569 non-null    float64
 8   concavity_mean           569 non-null    float64
 9   concave points_mean      569 non-null    float64
 10  symmetry_mean            569 non-null    float64
 11  fractal_dimension_mean   569 non-null    float64
 12  radius_se                569 non-null    float64
 13  texture_se               569 non-null    float64
 14  perimeter_se             569 non-null    float64
 15  area_se                  569 non-null    float64
 16  smoothness_se            569 non-null    float64
 17  compactness_se           569 non-null    float64
 18  concavity_se             569 non-null    float64
 19  concave points_se        569 non-null    float64
 20  symmetry_se              569 non-null    float64
 21  fractal_dimension_se     569 non-null    float64
 22  radius_worst             569 non-null    float64
 23  texture_worst            569 non-null    float64
 24  perimeter_worst          569 non-null    float64
 25  area_worst               569 non-null    float64
 26  smoothness_worst         569 non-null    float64
 27  compactness_worst        569 non-null    float64
 28  concavity_worst          569 non-null    float64
 29  concave points_worst     569 non-null    float64
 30  symmetry_worst           569 non-null    float64
 31  fractal_dimension_worst  569 non-null    float64
 32  Unnamed: 32              0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

In [6]:

```
df.describe()
```

Out[6]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.0
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.0
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.0
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.0
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.0
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.0
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.1
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.1

8 rows × 32 columns

In [7]:

```
df.columns
```

Out[7]:

```
Index(['id', 'diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
       'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mea
n',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_s
e',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_s
e',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst', 'Unnamed: 32'],
      dtype='object')
```

In [8]:

```
sb.pairplot(df)
```

Out[8]:

```
<seaborn.axisgrid.PairGrid at 0x20830bdf850>
```

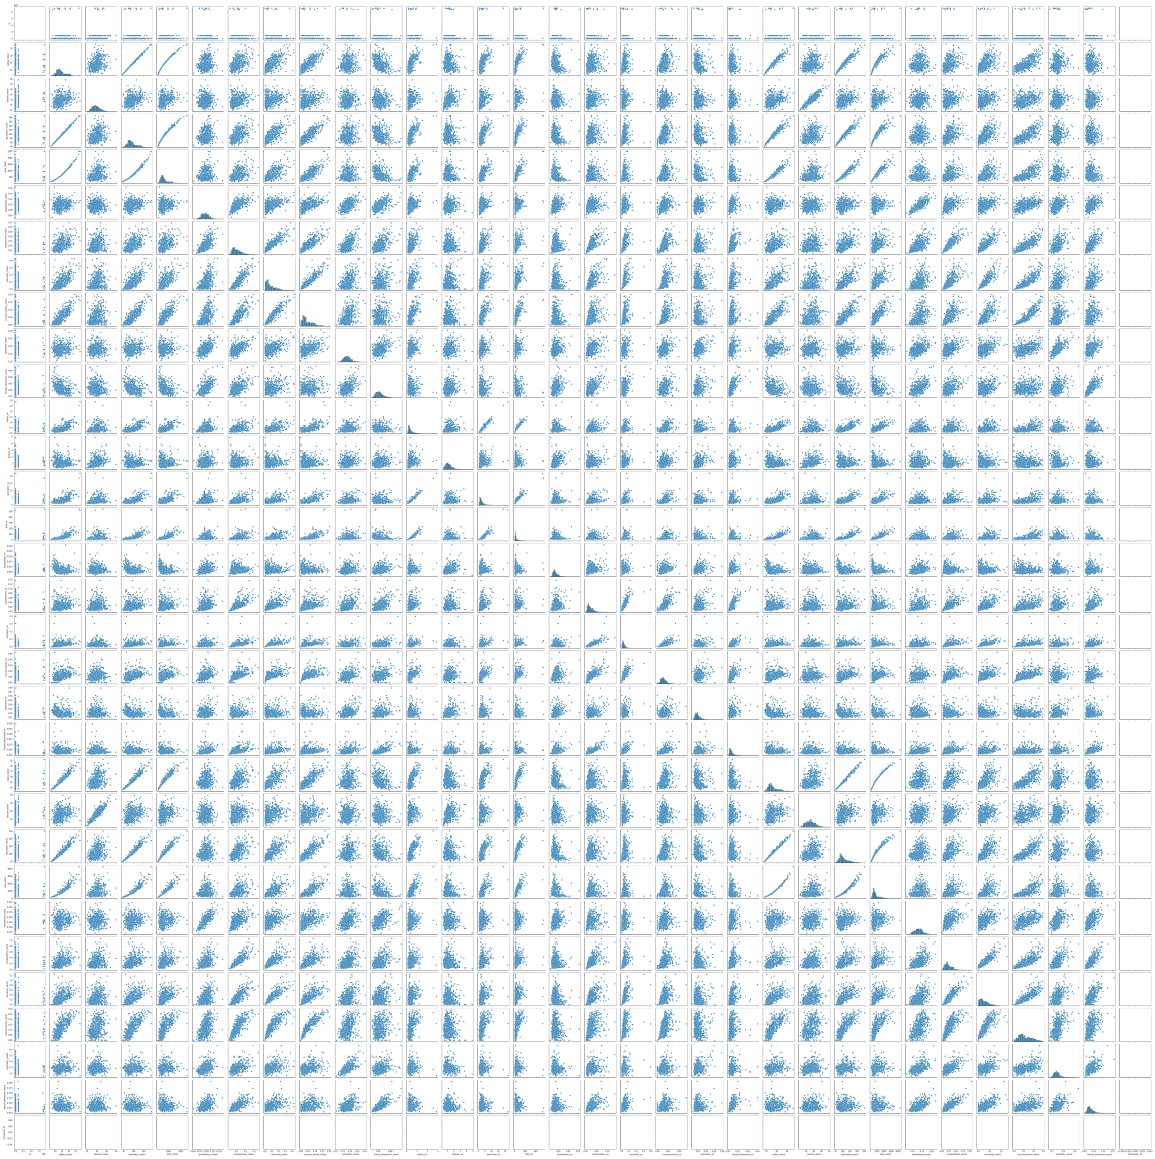

In [9]:

```
sb.distplot(df["radius_mean"])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[9]:

```
<AxesSubplot:xlabel='radius_mean', ylabel='Density'>
```

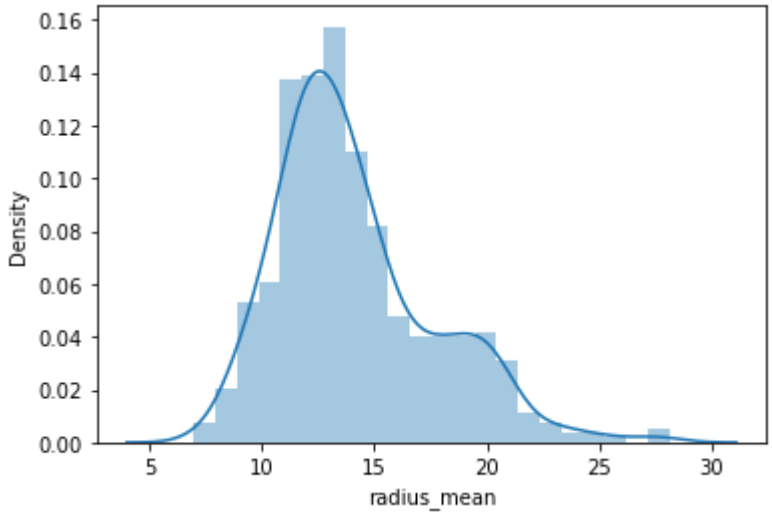

In [10]:

```
sb.distplot(df["radius_mean"])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[10]:

```
<AxesSubplot:xlabel='radius_mean', ylabel='Density'>
```

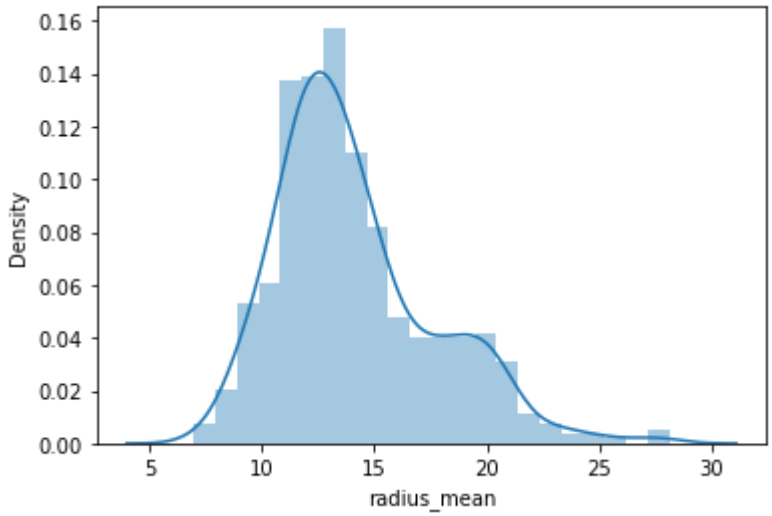

In [17]:

```
df1=df[['id','radius_mean', 'texture_mean', 'perimeter_mean',
       'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst']]
df1
```

Out[17]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_r
0	842302	17.99	10.38	122.80	1001.0	0.11840	0.2
1	842517	20.57	17.77	132.90	1326.0	0.08474	0.0
2	84300903	19.69	21.25	130.00	1203.0	0.10960	0.1
3	84348301	11.42	20.38	77.58	386.1	0.14250	0.2
4	84358402	20.29	14.34	135.10	1297.0	0.10030	0.1
...	...	...	...	...	...	...	
564	926424	21.56	22.39	142.00	1479.0	0.11100	0.1
565	926682	20.13	28.25	131.20	1261.0	0.09780	0.1
566	926954	16.60	28.08	108.30	858.1	0.08455	0.1

In [18]:

```
sb.heatmap(df1.corr())
```

Out[18]:

```
<AxesSubplot:>
```

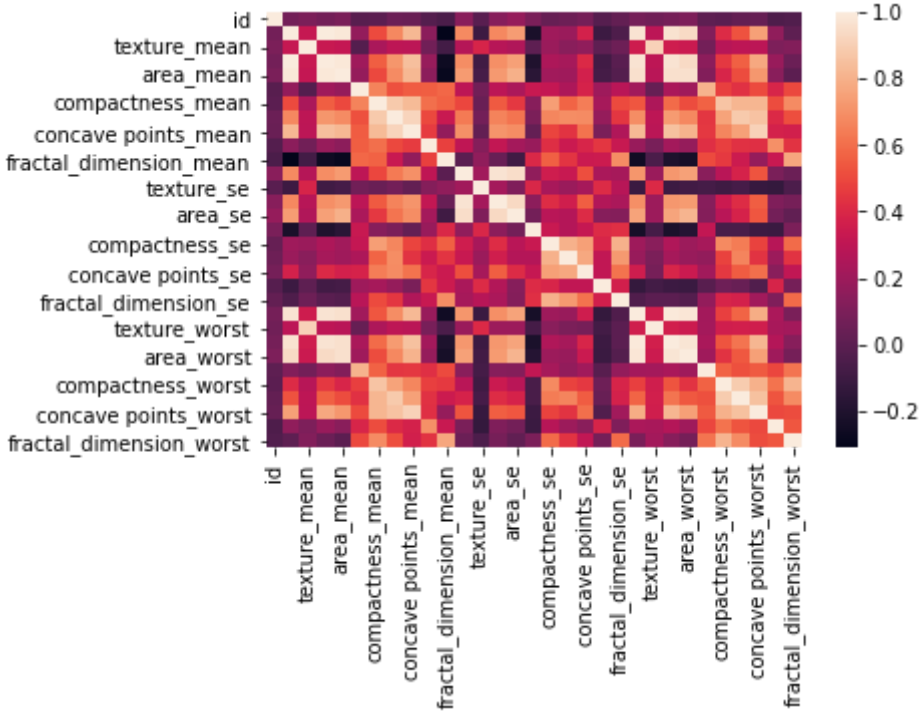

model building

In [19]:

```
x = df1[['id', 'radius_mean', 'texture_mean', 'perimeter_mean',
       'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst']]
y = df1['radius_mean']
```

In [20]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear regression

In [21]:

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(x_train,y_train)
```

Out[21]:

```
LinearRegression()
```

In [22]:

```
print(lr.intercept_)
```

```
-2.9868606254979113e-07
```

In [23]:

```
coef = pd.DataFrame(lr.coef_,x.columns,columns=['Co_efficient'])
coef
```

Out[23]:

	Co_efficient
id	-1.463460e-17
radius_mean	9.999999e-01
texture_mean	-7.577773e-09
perimeter_mean	1.848834e-08
area_mean	-3.039308e-10
smoothness_mean	-1.443566e-07
compactness_mean	-8.507271e-07
concavity_mean	7.441982e-08
concave points_mean	-4.808865e-07
symmetry_mean	5.414591e-07
fractal_dimension_mean	-1.930360e-07
radius_se	6.682006e-07
texture_se	-6.537640e-08
perimeter_se	-6.407581e-08
area_se	-1.254316e-09
smoothness_se	6.111042e-07
compactness_se	8.220824e-07
concavity_se	8.512614e-07
concave points_se	7.136502e-07
symmetry_se	-1.216300e-08
fractal_dimension_se	1.597606e-07
radius_worst	-9.694136e-08
texture_worst	9.343836e-09
perimeter_worst	8.912356e-09
area_worst	1.984561e-10
smoothness_worst	8.910018e-07
compactness_worst	-3.288526e-07
concavity_worst	-7.478191e-08
concave points_worst	-2.069475e-07
symmetry_worst	3.996945e-07
fractal_dimension_worst	-9.356701e-08

In [24]:

```
print(lr.score(x_test,y_test))
```

```
0.9999999999999998
```

In [25]:

```
prediction = lr.predict(x_test)
pp.scatter(y_test,prediction)
```

Out[25]:

```
<matplotlib.collections.PathCollection at 0x208691bf730>
```

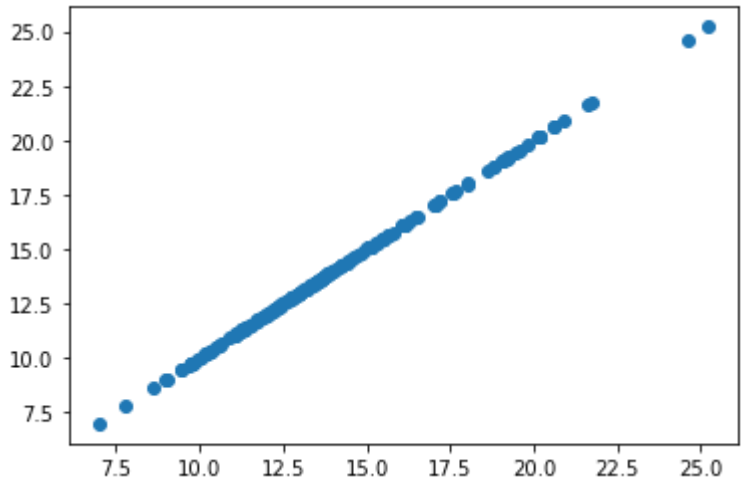

lasso and ridge regression

In [26]:

```
lr.score(x_test,y_test)
```

Out[26]:

```
0.9999999999999998
```

In [27]:

```
lr.score(x_train,y_train)
```

Out[27]:

```
0.9999999999999998
```

In [28]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [29]:

```
r = Ridge(alpha=10)
r.fit(x_train,y_train)
r.score(x_test,y_test)
r.score(x_train,y_train)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_ridge.py:
147: LinAlgWarning: Ill-conditioned matrix (rcond=7.67676e-19): result may
not be accurate.
  return linalg.solve(A, Xy, sym_pos=True,
```

Out[29]:

```
0.9996864978685623
```

In [30]:

```
l = Lasso(alpha=10)
l.fit(x_train,y_train)
l.score(x_test,y_test)
l.score(x_train,y_train)
```

Out[30]:

```
0.9740085529406495
```

elasticnet

In [31]:

```
from sklearn.linear_model import ElasticNet
e = ElasticNet()
e.fit(x_train,y_train)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinat
e_descent.py:530: ConvergenceWarning: Objective did not converge. You migh
t want to increase the number of iterations. Duality gap: 21.3713803495654
7, tolerance: 0.5225839749489949
  model = cd_fast.enet_coordinate_descent(
```

Out[31]:

```
ElasticNet()
```

In [32]:

```
print(e.coef_)
```

```
[-3.06107630e-11  0.00000000e+00  0.00000000e+00  7.62708499e-02
  5.18379514e-03 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00
 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00  0.00000000e+00
 -0.00000000e+00 -0.00000000e+00 -5.64214194e-03 -0.00000000e+00
 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00
 -0.00000000e+00  0.00000000e+00  0.00000000e+00  4.44155526e-03
 -2.04938933e-04 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00
 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00]
```

In [33]:

```
print(e.intercept_)
```

```
3.6494277052006296
```

In [34]:

```
predictions = e.predict(x_test)
predictions
```

Out[34]:

```
array([12.1714451 , 12.31129047, 13.44428327, 12.97869634, 13.31885489,
       11.54309161, 12.41047998, 19.89156585, 13.5690934 , 10.22992136,
       21.29111511, 12.44908191, 17.5595404 , 14.3068202 , 11.31430865,
       17.39258085, 13.69661883, 14.07461369, 16.12806409,  8.34225971,
       10.47868754, 12.64206486, 12.02499977, 15.51014042, 10.72660646,
       15.29502459, 13.17158501, 12.80156454, 20.15788721, 12.15420731,
       11.36331046, 13.03741273, 14.01544602, 14.42804333, 13.65113089,
       12.6760864 , 12.04639411, 11.40716704, 16.46995354, 11.86815295,
       12.5554015 , 12.50940673, 16.64718399, 13.52540169, 11.2433992 ,
       15.00851182, 13.35038681, 13.51885364, 11.99060778, 13.35591443,
       11.4105054 , 16.80067421, 14.61720893, 15.09852617, 18.56873895,
       20.94265846,  9.02346781, 10.03478903, 13.47166274,  9.86985112,
       13.65070864, 13.19682162, 13.67155183, 12.66955376, 10.98040211,
       12.7252163 , 14.07208511, 11.57924606, 13.15692964, 14.05639157,
       10.22892608, 11.33875948, 12.01664161, 19.41394745,  9.38156853,
       13.24440651, 13.64219128, 13.62770273, 12.3910594 , 12.78531875,
       15.22442262, 11.29385598, 15.70204115, 10.61484476, 11.69136889,
       15.00172765, 13.76537966, 18.51755687, 17.86838968, 12.91703413,
       11.73683728,  9.60578052, 19.96779455, 12.23501444, 19.80025628,
       18.58616946, 12.94196094, 14.79738919, 13.14871584, 18.95146111,
       19.41215196, 10.39918676, 13.37838825,  9.86662932, 14.38799131,
       15.78219392, 14.21399194, 12.81233352, 14.79110446, 11.50441004,
       13.7686483 , 11.67753635, 25.39565488, 10.67190163, 17.82801352,
        9.78038232, 11.92930399, 13.35521314, 11.44469197, 16.06494326,
       26.20294879, 15.94883243, 15.66967344, 11.42102545, 17.11422948,
       11.98793257, 12.14657883, 18.55082671, 20.83014585, 10.98682228,
       11.17460053, 13.78830367, 10.39659813, 12.25753722, 12.41333744,
       11.99403369, 18.9307201 , 11.38443221, 14.90544098, 10.02830829,
        9.91610514,  7.86364134, 16.11677138, 13.14243084, 12.17069808,
       13.50533266, 12.95442689, 13.86864655, 12.04839452,  9.63894143,
       20.85269807, 14.79333959, 21.83702304, 17.37456292, 16.24846695,
       20.3833899 , 17.10685127, 19.34678021, 15.39594338, 11.16805829,
       15.06894868, 14.34003422, 10.03068179, 15.63151999, 12.80258082,
       14.80399087, 13.41954054, 22.10972808, 10.23360876, 17.04682446,
       10.49825453])
```

In [35]:

```
print(e.score(x_test,y_test))
```

```
0.9948641578483264
```

In [36]:

```
from sklearn import metrics
```

mean absolute error

In [37]:

```
print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,predictions))
```

```
Mean Absolute Error: 0.1773948243706064
```

mean squared error

In [38]:

```
print("Mean Squared Error:", metrics.mean_squared_error(y_test,predictions))
```

```
Mean Squared Error: 0.054725972652547986
```

root mean squared error

In [39]:

```
print("Root Mean Squared Error",np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

```
Root Mean Squared Error 0.23393583020253222
```

In [ ]:

```

```